

Text classification on real data.

This notebook based on my laboratory work. Only minor adjustments have been made.

Getting data

The reviews were taken from the marketplace <https://www.wildberries.ru>. A parser was developed that uses the `requests` and `pydantic` libraries. The data on the site page is loaded dynamically using requests to the API (for example, to server with reviews, to server with product catalogs, etc.). Using the `requests` library functions, requests are made to wildberries, then the response is parsed according to the specified data models and their attributes based on the `pydantic` class.

- `models.py` stores data models;
- `parser.py` implements the parser class;

The algorithm of the parser:

- Initializes the parser class with links to different product from the marketplace
- The product ID is parsed from the link using the `re` regular expression library
- A request is made to wildberries, with the parameter obtained in the previous step, and the ID of the seller is extracted from the request, whose goods will then be parsed in the cycle
- A csv file is created with the columns `name`, `brand`, `price`, `sale`, `price_with_sale`, `number of photos in the card - pics`, `pros`, `cons`, `availability of photos in the review - hasPhoto`, `useful`, `useless`, `text`, `rating` - `target`
- As long as the server returns a non-empty json response, we make a request to it, receiving a list of the seller's products page by page, with the fields above
- Receiving a list of products, we make a request to the server with feedback - reviews for one product from the list, process the received data and put them in a csv file

As links to products, I randomly selected products from the site from different categories of different sellers to make the sample representative.

Import necessary libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import nltk
nltk.download("stopwords")
from nltk.corpus import stopwords
from wordcloud import WordCloud, STOPWORDS
from pymorphy3 import MorphAnalyzer

from sklearn.preprocessing import StandardScaler
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import train_test_split
from imblearn.under_sampling import RandomUnderSampler
from sklearn.metrics import classification_report, accuracy_score

%matplotlib inline
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\ibas1\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Cleaning data and descriptive statistics

Read data, rename columns from Russian to English

```
In [2]: df = pd.read_csv("../data/dataset.csv")
df.rename(columns={
    "название": "name",
    "бренд": "brand",
    "цена": "price",
    "скидка в %": "sale",
    "цена со скидкой": "price_with_sale",
    "кол-во фото в карточке": "pics",
    "плюсы": "pros",
    "минусы": "cons",
    "наличие фото в отзыве": "hasPhoto",
    "полезно": "useful",
    "неполезно": "useless",
    "текст отзыва": "text",
    'рейтинг': 'target'},
    inplace=True)
df.head()
```

```
C:\Users\ibas1\AppData\Local\Temp\ipykernel_17344\2546511965.py:1: DtypeWarning: Columns (8) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv("../data/dataset.csv")
```

Out[2]:

	name	brand	price	sale	price_with_sale	pics	pros	cons	hasPhoto	useful	use
0	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	204.0	5	NaN	NaN	False	0	
1	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	204.0	5	NaN	NaN	False	0	
2	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	204.0	5	NaN	NaN	False	0	
3	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	204.0	5	NaN	NaN	False	0	
4	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	204.0	5	NaN	NaN	False	0	

Drop duplicates and view shape

```
In [3]: df.drop_duplicates(inplace=True)  
df.shape
```

Out[3]: (381859, 13)

Main information about dataset

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 381859 entries, 0 to 461663
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                   381859 non-null object
1   brand                  381857 non-null object
2   price                  381859 non-null float64
3   sale                   381859 non-null int64
4   price_with_sale        381859 non-null float64
5   pics                   381859 non-null int64
6   pros                   0 non-null      float64
7   cons                   0 non-null      float64
8   hasPhoto               381033 non-null object
9   useful                 381859 non-null int64
10  useless                 381859 non-null int64
11  text                   380199 non-null object
12  target                 381859 non-null int64
dtypes: float64(4), int64(5), object(4)
memory usage: 40.8+ MB
```

Change type of hasPhoto column

```
In [5]: df['hasPhoto'] = df['hasPhoto'].astype(bool)
```

Descriptive statistics

```
In [6]: df.describe()
```

```
Out[6]:
```

	price	sale	price_with_sale	pics	pros	cons	useful
count	381859.000000	381859.000000	381859.000000	381859.000000	0.0	0.0	381859.000000
mean	8620.201093	58.793353	3335.807565	6.952530	NaN	NaN	1.337533
std	17832.819749	20.181132	6753.828420	4.154662	NaN	NaN	4.079684
min	99.000000	0.000000	85.000000	1.000000	NaN	NaN	0.000000
25%	1500.000000	59.000000	465.000000	4.000000	NaN	NaN	0.000000
50%	3200.000000	66.000000	1288.000000	6.000000	NaN	NaN	0.000000
75%	7030.000000	70.000000	3146.000000	9.000000	NaN	NaN	1.000000
max	629990.000000	93.000000	239396.000000	26.000000	NaN	NaN	332.000000

Descriptive statistics of object columns

```
In [7]: df.describe(include=['O'])
```

Out[7]:


	name	brand	text
count	381859	381857	380199
unique	2012	102	201389
top	Ювелирные серьги женские из серебра 925 SOKOLOV все отлично		
freq	9322	178255	716

Amount of missed values

In [8]: `df.isna().sum()`

```
Out[8]: name          0
brand          2
price          0
sale           0
price_with_sale 0
pics           0
pros          381859
cons          381859
hasPhoto       0
useful         0
useless        0
text           1660
target         0
dtype: int64
```

In [9]: `sns.heatmap(df.isna(), yticklabels=False, cbar=False, cmap='Blues');`



pros and cons columns contain only empty values, remove them. Also we delete all the remaining missed values because we have a lot of data


```
In [10]: df.drop(columns=['pros', 'cons'], inplace=True)
df.dropna(inplace=True)
df.drop_duplicates(subset=['text'], inplace=True)
df.shape
```

```
Out[10]: (201388, 11)
```

Correlation matrix

```
In [11]: numeric_cols = df.select_dtypes(include=[np.number]).columns
df_ = df[numeric_cols]
corr = df_.corr()
mask = np.zeros_like(corr, dtype=np.bool_)
mask[np.triu_indices_from(mask)] = True

plt.figure(figsize = (10,7))
plt.title('Correlation matrix')
sns.heatmap(corr,
            mask=mask,
            annot=True,
            fmt='.2g',
            linewidths=2);
```







price_with_sale and price correlate with the value 0.97. Delete price_with_sale


```
In [12]: df.drop(columns=['price_with_sale'], inplace=True)
```

Look at the outliers in numerical data using the boxplot

```
In [13]: numeric_columns = df.select_dtypes(include=np.number).columns.tolist()
for col in numeric_columns:
    plt.figure(figsize=(4,4))
    sns.boxplot(data=df[col])
    plt.title(f'{col}')
    plt.ylabel('Принимаемые значения');
```







In my opinion, we can remove a little bit of outliers in the `useful` and `useless` columns,


```
In [14]: Q1 = df['price'].quantile(0.25)
Q3 = df['price'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

df[(df['price'] >= lower_bound) & (df['price'] <= upper_bound)].shape
```


Out[14]: (174487, 10)

```
In [15]: plt.figure(figsize=(4,4))
sns.boxplot(data=df[(df['price'] >= lower_bound) & (df['price'] <= upper_bound)]['p
plt.title(f'Price')
plt.ylabel('Values');
```



We will remove ~15% of the emissions from the `price` column. There is a lot of data, it is not critical.

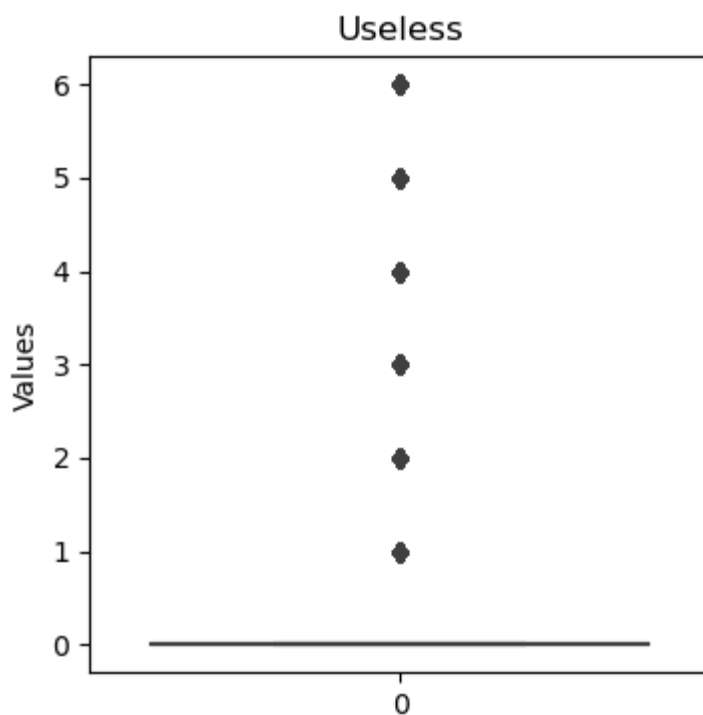
```
In [16]: plt.figure(figsize=(4,4))
sns.boxplot(data=df['useful'][df['useful'] < df['useful'].quantile(0.99)])
plt.title(f'Usefull')
plt.ylabel('Values');
```



```
In [17]: df['useful'].value_counts()
```

```
Out[17]: useful
0      112014
1       40038
2       17313
3        8874
4        5513
...
65         1
108        1
110        1
114        1
90         1
Name: count, Length: 127, dtype: int64
```

```
In [18]: plt.figure(figsize=(4,4))
sns.boxplot(data=df['useless'][df['useless'] < df['useless'].quantile(0.99)])
plt.title(f'Useless')
plt.ylabel('Values');
```



```
In [19]: df['useless'].value_counts()
```

```
Out[19]: useless
0      160961
1       24610
2        7409
3        3169
4        1610
...
103         1
63          1
82          1
47          1
72          1
Name: count, Length: 82, dtype: int64
```

Delete 1% of useless and useful .

```
In [20]: df = df[df['price'] < df['price'].quantile(0.85)]
df = df[df['useless'] < df['useless'].quantile(0.99)]
df = df[df['useful'] < df['useful'].quantile(0.99)]
df.shape
```

```
Out[20]: (166779, 10)
```

Look at the basic information of the cleaned dataset

```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 166779 entries, 0 to 461608
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   name        166779 non-null object  
 1   brand       166779 non-null object  
 2   price       166779 non-null float64 
 3   sale        166779 non-null int64  
 4   pics        166779 non-null int64  
 5   hasPhoto    166779 non-null bool   
 6   useful      166779 non-null int64  
 7   useless     166779 non-null int64  
 8   text        166779 non-null object  
 9   target      166779 non-null int64  
dtypes: bool(1), float64(1), int64(5), object(3)
memory usage: 12.9+ MB
```

EDA

In [22]: `df.head()`

Out[22]:

	name	brand	price	sale	pics	hasPhoto	useful	useless	text	target
0	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	5	False	0	1	отличные кусачки заточены хорошо со своей зада...	5
1	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	5	False	0	1	спасибо за качественный товар буду рекомендова...	5
2	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	5	False	0	0	щипчики хорошо стригут ногти все отлично	5
3	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	5	False	0	0	хороший набор пришло все целое	5
4	Клиппер маникюрный, кусачки для ногтей	Zebo Professional	499.0	59	5	False	0	0	получили кусочки пришли быстро и хорошо упаков...	5


Target rating

In [23]: `df['target'].value_counts()`


```
Out[23]: target
5      134566
4       12441
1        8079
3        7812
2        3881
Name: count, dtype: int64
```

```
In [24]: df["target"].value_counts().plot(
        kind='pie',
        title='Rating',
        figsize=(7, 7),
        autopct='%0.2f%%')

plt.ylabel('');
```



```
In [25]: ax = sns.countplot(x='target', data=df)
ax.bar_label(ax.containers[0]);
```





Classes are not balanced - there are too many "5" and too few other grades, they are 4 times less. Since there is a lot of data, we will do undersampling. Each class will have 3800 - this is enough for our model to have good generalization capability.

In addition, it can be concluded that in 4 out of 5 cases, the buyer is satisfied with the purchased product, which is a good indicator.

```
In [26]: avg_rate_by_brand = pd.DataFrame(df.groupby('brand')['target'].mean())\
                                                .reset_index()\
                                                .sort_values(by='target', ascending=False)

plt.figure(figsize=(5, 15))
sns.barplot(x='target', y='brand', data=avg_rate_by_brand, orient='h')
plt.xlabel('Average rating')
plt.ylabel('Brand')
plt.title('Average rating by brands')
plt.show()
```








Most brands keep the quality - the average rating is 4 or higher, but there are also unscrupulous ones who sell goods of poor quality, which have a rating of 3 or lower, but there are only 6 such manufacturers.

```
In [27]: avg_price_by_brand = pd.DataFrame(df.groupby('brand')['price'].mean())\
        .reset_index()\
        .sort_values(by='price', ascending=False)

plt.figure(figsize=(5, 15))
sns.barplot(x='price', y='brand', data=avg_price_by_brand, orient='h')
plt.xlabel('Average price')
plt.ylabel('Brand')
plt.title('Average price by brand')
plt.show()
```



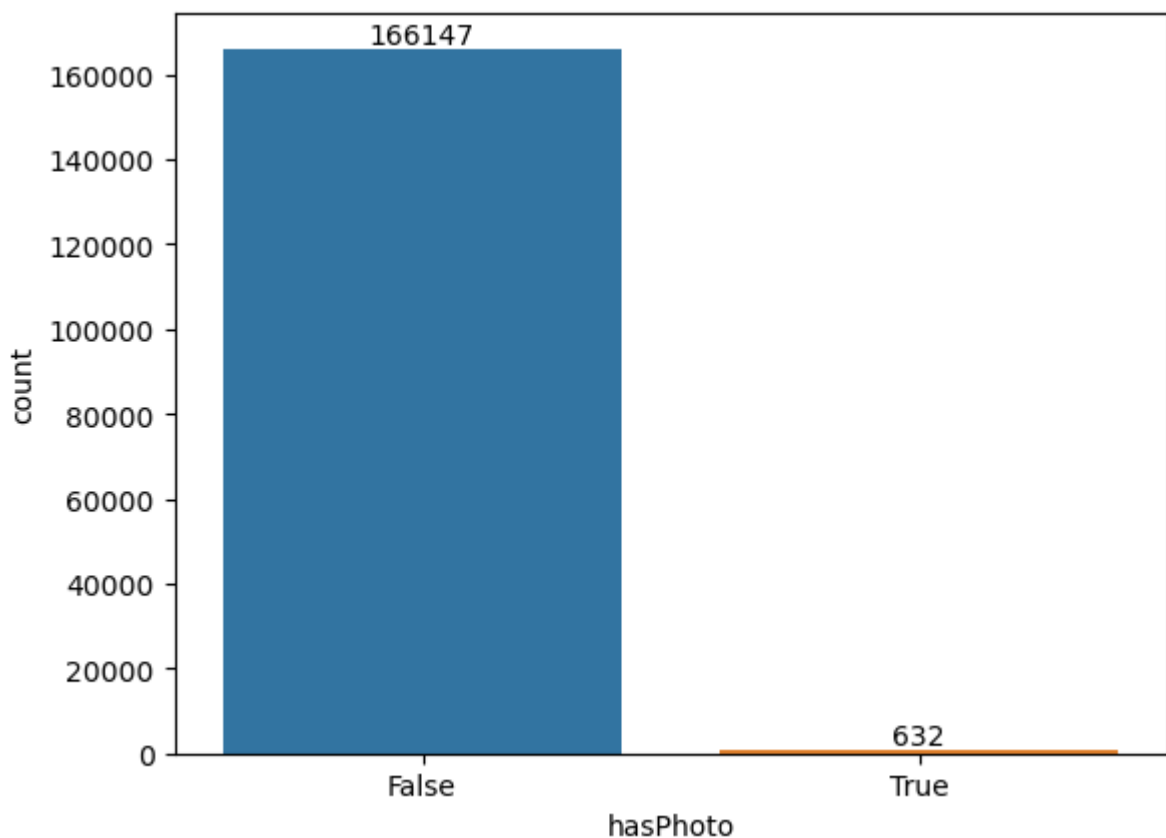


Most of the products of various brands are sold at an average price of up to ~5000 rubles.

It can also be noted that the closing three brands according to the average rating have a price for goods of 4-6 thousand rubles, i.e. for relatively big money, buyers receive a terrible product.

Column hasPhoto

```
In [28]: ax = sns.countplot(x='hasPhoto', data=df)
ax.bar_label(ax.containers[0]);
```




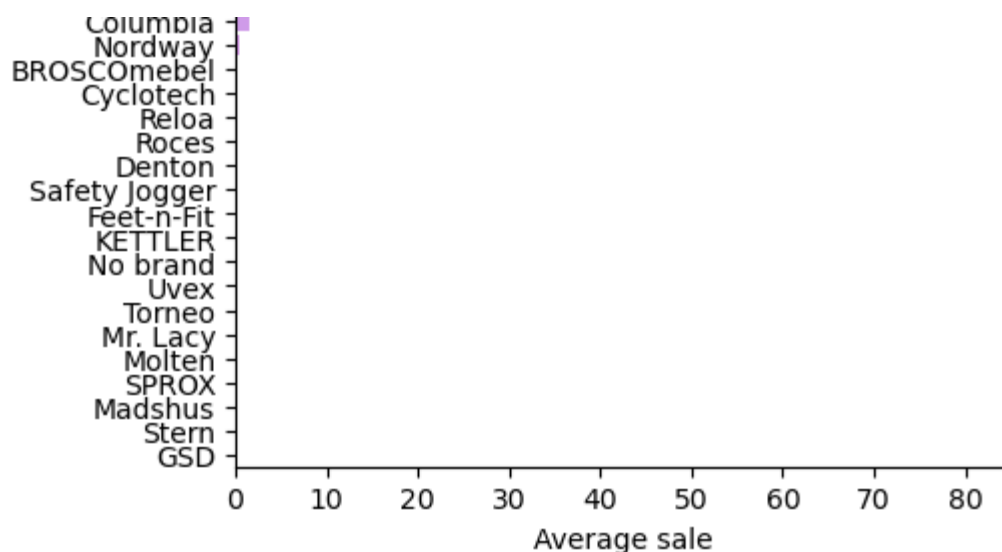
Conclusion: the vast majority of users prefer to leave reviews without photos.

Column sale

```
In [29]: avg_sale_by_brand = pd.DataFrame(df.groupby('brand')['sale'].mean())\
                                                .reset_index()\
                                                .sort_values(by='sale', ascending=False)

plt.figure(figsize=(5, 15))
sns.barplot(x='sale', y='brand', data=avg_sale_by_brand, orient='h')
plt.xlabel('Average sale')
plt.ylabel('Brand')
plt.title('Average sale by brand')
plt.show()
```

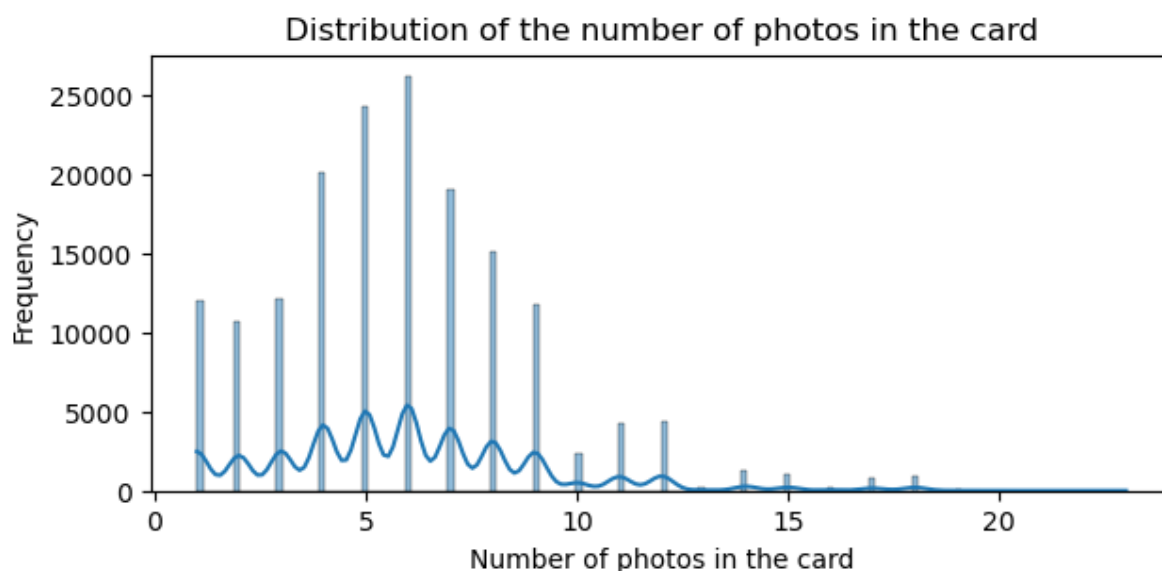




About half of the brands have a discount of more than 40%, which is undoubtedly a cheat, since this is done to promote product cards to the top of the list and create the effect of acquiring benefits for the buyer. About 20% of brands do not have a discount.


Column `pics` - number of images in product card

```
In [30]: plt.figure(figsize=(7,3))
sns.histplot(x=df['pics'], kde=True)
plt.title('Distribution of the number of photos in the card')
plt.xlabel('Number of photos in the card')
plt.ylabel('Frequency');
```




Most often, sellers take from 1 to 9 photos in the card. Let's look at the product rating depending on the number of photos in the product card

```
In [31]: plt.figure(figsize=(15,5))
ax = sns.countplot(x='pics', data=df, hue='target');
```



There is no special dependence

```
In [32]: ax = sns.countplot(x='target', data=df)
ax.bar_label(ax.containers[0]);
```



As I said earlier, the classes are unbalanced, and due to the fact that there is a lot of data, we will do undersampling

Text processing

Word cloud


```
In [33]: # get list of words from string
def str_corpus(corpus):
    str_corpus = ''
    for i in corpus:
        str_corpus += ' ' + i
    str_corpus = str_corpus.strip()
    return str_corpus

# get all words in corpus
def get_corpus(data):
    corpus = []
    for phrase in data:
        for word in phrase.split():
            corpus.append(word)
    return corpus

# getting word clord
def get_wordCloud(corpus):
    wordCloud = WordCloud(background_color='white',
                           stopwords=STOPWORDS,
                           width=3000,
                           height=2500,
                           max_words=200,
                           random_state=42
                           ).generate(str_corpus(corpus))


    return wordCloud

corpus = get_corpus(df['text'].values)
procWordCloud = get_wordCloud(corpus)

fig = plt.figure(figsize=(20, 8))
plt.subplot(1, 2, 1)
plt.imshow(procWordCloud)
plt.axis('off')
plt.subplot(1, 2, 1);
```



```
fig = plt.figure(figsize=(20, 8))
plt.subplot(1, 2, 1)
plt.imshow(procWordCloud)
plt.axis('off')
plt.subplot(1, 2, 1);
```



```
In [38]: X_resampled.reset_index(drop=True, inplace=True)
          y_resampled.reset_index(drop=True, inplace=True)

          df = pd.concat([X_resampled, y_resampled], axis=1)
```

```
corpus = get_corpus(df['text'].values)
procWordCloud = get_wordCloud(corpus)
```

```
fig = plt.figure(figsize=(20, 8))
plt.subplot(1, 2, 1)
plt.imshow(procWordCloud)
plt.axis('off')
plt.subplot(1, 2, 1);
```



Due to the fact that more than half of the dataset consists of reviews of Sokolov(jewelry company) products, there are many words related to the subject of jewelry (for example, silver, earrings, ring, etc.)

Processing of numerical and categorical features

Standardize numerical features, encode categorical features

```
numeric columns
```

```
['price', 'sale', 'pics', 'useful', 'useless', 'target']
```

```
In [41]: scaler = StandardScaler()
df[numeric_columns[:-1]] = scaler.fit_transform(df[numeric_columns[:-1]])
```

The history saving thread hit an unexpected error (OperationalError('database is locked')).History will not be written to the database.

```
In [42]: df
```

Out[42]:

	name	brand	price	sale	pics	hasPhoto	useful	useless	
0	Шлепанцы Mono	FILA	-0.144880	-2.472030	-0.618441	False	-0.030635	0.600217	с от
1	Кольцо из серебра	SOKOLOV	0.178815	0.401142	-0.925706	False	0.518779	-0.542338	т
2	Ювелирный каучуковый шнурок с замком из серебр...	SOKOLOV	0.015126	0.575274	-0.925706	False	1.617607	0.600217	па:
3	Ювелирные серьги кольца серебро 925	SOKOLOV	-0.230407	0.314076	-0.618441	False	-0.580049	-0.542338	т
4	Женские серьги пусеты гвоздики из золота 585 п...	SOKOLOV	2.715995	0.357609	0.610619	False	-0.580049	1.742772	по
...	
19400	Астильба Арендса Mix	Садовые Растения	-0.962916	0.749406	-1.540236	False	-0.030635	-0.542338	
19401	Кольцо из серебра с фианитом	SOKOLOV	-0.312252	0.314076	-0.311176	False	-0.580049	-0.542338	
19402	Женское кольцо на помолвку из серебра 925	SOKOLOV	0.015126	0.314076	0.917884	False	-0.580049	-0.542338	т а
19403	Ювелирная подвеска кулон на шею серебро 925	SOKOLOV	-0.475941	0.314076	0.610619	False	2.167021	0.600217	пс

	name	brand	price	sale	pics	hasPhoto	useful	useless
19404	Шлепанцы	FILA	-0.144880	-2.472030	-0.618441	False	-0.580049	-0.542338

19405 rows × 10 columns

Encode categorical features

```
In [43]: df['brand'].value_counts()
```

```
Out[43]: brand
SOKOLOV      9423
Zebo Professional  2060
Садовые Растения  1413
Demix        911
Permegi       650
...
Mountain Hardwear    1
HEAD                 1
Denton               1
Street Surfing       1
QREDO                1
Name: count, Length: 86, dtype: int64
```

```
In [44]: df['name'].value_counts()
```

```
Out[44]: name
Ювелирные серьги женские из серебра 925      711
Кольцо из серебра с фианитами                667
Цепочка на шею из серебра 925                423
Ювелирные серьги пусеты-гвоздики из серебра 925  394
Ювелирная цепочка на шею серебро 925          391
...
Куртка софтшелл женская                      1
Ролики JUNIOR GIRL                          1
Перчатки сноубордические KAILA               1
Виноград девичий Yellow Wall                 1
Куртка утепленная для девочек                 1
Name: count, Length: 988, dtype: int64
```

Encode brand by OneHotEncoder and remove its column

```
In [45]: brand = pd.get_dummies(df["brand"], drop_first=True, dtype=int)
df.drop(columns=["name", "brand"], inplace=True)
df = pd.concat([df, brand], axis=1)
df.head()
```


Out[45]:

	price	sale	pics	hasPhoto	useful	useless	text	target	Alpashop
0	-0.144880	-2.472030	-0.618441	False	-0.030635	0.600217	пришли абсолютно схожи оригиналом фото китайск...	1	0
1	0.178815	0.401142	-0.925706	False	0.518779	-0.542338	кольцо пришло согнутое думаю продавцу нужно пе...	1	0
2	0.015126	0.575274	-0.925706	False	1.617607	0.600217	шнурок вместо сильно разочаровал ваш товар хот...	1	0
3	-0.230407	0.314076	-0.618441	False	-0.580049	-0.542338	серьги еле надела это кошмар	1	0
4	2.715995	0.357609	0.610619	False	-0.580049	1.742772	понравилась застежка	1	0

5 rows × 93 columns

Encode hasPhoto by OneHotEncoder

```
In [46]: has_photo = pd.get_dummies(df["hasPhoto"], drop_first=True, dtype=int)
df.drop(columns=["hasPhoto"], inplace=True)
df = pd.concat([df, has_photo], axis=1)
df.head()
```


Out[46]:

	price	sale	pics	useful	useless	text	target	Alpashop	BASEUS
0	-0.144880	-2.472030	-0.618441	-0.030635	0.600217	пришли абсолютно схожи оригиналом фото китайск...	1	0	0
1	0.178815	0.401142	-0.925706	0.518779	-0.542338	кольцо пришло согнутое думаю продавцу нужно пе...	1	0	0
2	0.015126	0.575274	-0.925706	1.617607	0.600217	шнурок вместо сильно разочаровал ваш товар хот...	1	0	0
3	-0.230407	0.314076	-0.618441	-0.580049	-0.542338	серьги еле надела это кошмар	1	0	0
4	2.715995	0.357609	0.610619	-0.580049	1.742772	понравилась застежка	1	0	0

5 rows × 93 columns

Text processing

```
In [47]: morph = MorphAnalyzer()
def lemmatize(doc):
    tokens = []
    for token in doc.split():
        if token:
            token = token.strip()
            token = morph.normal_forms(token)[0]
            tokens.append(token)
    if len(tokens) > 0:
        return " ".join(tokens)
    return None

df["text"] = df["text"].apply(lemmatize)
df.head()
```

Out[47]:

	price	sale	pics	useful	useless	text	target	Alpashop	BASEUS
0	-0.144880	-2.472030	-0.618441	-0.030635	0.600217	прислать абсолютно схожий оригинал фото китайс...	1	0	0
1	0.178815	0.401142	-0.925706	0.518779	-0.542338	кольцо прийти согнутый думать продавец нужно п...	1	0	0
2	0.015126	0.575274	-0.925706	1.617607	0.600217	шнурок вместо сильно разочаровать ваш товар хо...	1	0	0
3	-0.230407	0.314076	-0.618441	-0.580049	-0.542338	серьга еле надеть это кошмар	1	0	0
4	2.715995	0.357609	0.610619	-0.580049	1.742772	понравиться застёжка	1	0	0

5 rows × 93 columns

In [48]: `df["text"].isna().sum()`

Out[48]: 6

In [49]: `df.dropna(inplace=True)`In [50]: `df.head()`

Out[50]:

	price	sale	pics	useful	useless	text	target	Alpashop	BASEUS
0	-0.144880	-2.472030	-0.618441	-0.030635	0.600217	прислать абсолютно схожий оригинал фото китайс...	1	0	0
1	0.178815	0.401142	-0.925706	0.518779	-0.542338	кольцо прийти согнутый думать продавец нужно п...	1	0	0
2	0.015126	0.575274	-0.925706	1.617607	0.600217	шнурок вместо сильно разочаровать ваш товар хо...	1	0	0
3	-0.230407	0.314076	-0.618441	-0.580049	-0.542338	серьга еле надеть это кошмар	1	0	0
4	2.715995	0.357609	0.610619	-0.580049	1.742772	понравиться застёжка	1	0	0

5 rows × 93 columns

```
In [51]: corpus = df["text"].to_list()

vectorizer = TfidfVectorizer()

X_tf_idf = vectorizer.fit_transform(corpus)

X_tf_idf
```

```
Out[51]: <19399x12312 sparse matrix of type '<class 'numpy.float64'>'
         with 189726 stored elements in Compressed Sparse Row format>
```

```
In [52]: tfidf_df = pd.DataFrame(X_tf_idf.toarray(), columns=vectorizer.get_feature_names_out())
```

```
In [53]: df.reset_index(drop=True, inplace=True)
         tfidf_df.reset_index(drop=True, inplace=True)

         df = pd.concat([df, tfidf_df], axis=1)
```

```
In [54]: df.drop(columns=["text"], inplace=True)
```

```
In [55]: df.columns = df.columns.astype(str)
```

```
In [56]: X = df.drop(columns=["target"])
         y = df["target"]
```

```
In [57]: y.shape, X.shape, df.shape
```

```
Out[57]: ((19399,), (19399, 12403), (19399, 12404))
```

Trainig model

```
In [58]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

KNN

```
In [59]: knnclf = KNeighborsClassifier()
```

Train

```
In [60]: %%time
knnclf.fit(X_train, y_train)
```

```
CPU times: total: 2.02 s
Wall time: 7.08 s
```

```
Out[60]: ▼ KNeighborsClassifier
KNeighborsClassifier()
```

Test

```
In [61]: %%time
y_pred = knnclf.predict(X_test)
print(f'accuracy {accuracy_score(y_pred, y_test):.2}')
print(classification_report(y_test, y_pred))
```

```
accuracy 0.33
```

	precision	recall	f1-score	support
1	0.34	0.46	0.39	767
2	0.26	0.27	0.26	773
3	0.24	0.19	0.21	792
4	0.30	0.25	0.27	785
5	0.48	0.47	0.47	763
accuracy			0.33	3880
macro avg	0.32	0.33	0.32	3880
weighted avg	0.32	0.33	0.32	3880

```
CPU times: total: 1min 46s
Wall time: 11.7 s
```

Random Forest

```
In [62]: rfc = RandomForestClassifier()
```

Train

```
In [63]: %%time
rfc.fit(X_train, y_train)
```

```
CPU times: total: 1min 11s
Wall time: 1min 31s
```

```
Out[63]: ▼ RandomForestClassifier
RandomForestClassifier()
```

Test

```
In [64]: %%time
y_pred = rfc.predict(X_test)
print(f'accuracy {accuracy_score(y_pred, y_test):.2}')
print(classification_report(y_test, y_pred))
```

```
accuracy 0.45
```

	precision	recall	f1-score	support
1	0.44	0.63	0.52	767
2	0.32	0.26	0.29	773
3	0.31	0.21	0.25	792
4	0.42	0.36	0.39	785
5	0.65	0.79	0.71	763
accuracy			0.45	3880
macro avg	0.43	0.45	0.43	3880
weighted avg	0.43	0.45	0.43	3880

```
CPU times: total: 422 ms
Wall time: 640 ms
```

AdaBoost

```
In [65]: abc = AdaBoostClassifier()
```

Train

```
In [66]: %%time
abc.fit(X_train, y_train)
```

```
CPU times: total: 56 s
Wall time: 1min 15s
```

```
Out[66]: ▼ AdaBoostClassifier
AdaBoostClassifier()
```

Test

```
In [67]: %%time
y_pred = abc.predict(X_test)
print(f'accuracy {accuracy_score(y_pred, y_test):.2f}')
print(classification_report(y_test, y_pred))
```

```
accuracy 0.4
```

	precision	recall	f1-score	support
1	0.43	0.55	0.49	767
2	0.28	0.15	0.19	773
3	0.28	0.30	0.29	792
4	0.35	0.38	0.36	785
5	0.62	0.64	0.63	763
accuracy			0.40	3880
macro avg	0.39	0.40	0.39	3880
weighted avg	0.39	0.40	0.39	3880

```
CPU times: total: 4.5 s
Wall time: 6.32 s
```

MLP

```
In [68]: mlpc = MLPClassifier()
```

Train

```
In [69]: %%time
mlpc.fit(X_train, y_train)
```

```
CPU times: total: 3h 57min 30s
Wall time: 49min 5s
```

```
C:\Users\ibas1\anaconda3\Lib\site-packages\sklearn\normalization\_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

```
warnings.warn(
Out[69]: ▼ MLPClassifier
MLPClassifier()
```

Test

```
In [70]: %%time
y_pred = mlpc.predict(X_test)
print(f'accuracy {accuracy_score(y_pred, y_test):.2f}')
print(classification_report(y_test, y_pred))
```

accuracy	0.41				
		precision	recall	f1-score	support
	1	0.45	0.47	0.46	767
	2	0.32	0.32	0.32	773
	3	0.26	0.26	0.26	792
	4	0.38	0.35	0.36	785
	5	0.66	0.69	0.68	763
accuracy				0.41	3880
macro avg		0.41	0.42	0.41	3880
weighted avg		0.41	0.41	0.41	3880

CPU times: total: 1.44 s

Wall time: 570 ms

```
In [71]: column0 = ['KNN', 'Random Forest', 'AdaBoost', 'MLP']
column1 = ["8.49 s", '1min 23s', '1 min 10s', '36 min 18s']
column2 = ["658 ms", '1min 22s', '1min 5s', '36min 17s']
column3 = ['7.82 s', '605 ms', '5.12 s', '406 ms']
column4 = [0.33, 0.45, 0.4, 0.42]

data = {'Model': column0,
        'Run time': column1,
        'Traing time': column2,
        'Test time': column3,
        'Accuracy': column4}
table = pd.DataFrame(data)
table.set_index('Model', inplace=True)
table
```

```
Out[71]:
```

	Run time	Traing time	Test time	Accuracy
Model				
KNN	8.49 s	658 ms	7.82 s	0.33
Random Forest	1min 23s	1min 22s	605 ms	0.45
AdaBoost	1 min 10s	1min 5s	5.12 s	0.40
MLP	36 min 18s	36min 17s	406 ms	0.42

Conclusion

The overall quality of the models is average - each of them is very close to random guessing the rating, with the exception of reviews with a rating of "5". The random forest and multilayer perceptron models distinguish five stars texts much better than the other - with 0.63 and 0.7 f1-score. In my opinion, this is due to four factors - the comments are not long enough for the model to extract the desired message; you need to use word embeddings instead of tf-idf vectorization; you need to use more complex algorithms, for example, deep learning with rnn or with attention layers; the authors' comments are biased and vary;

I said average quality because even modern BERT models do not give an accuracy of more than 0.6 due to the [article](#)

Table 2: SST-5 Results

Model	Training Time (per epoch)	Best Test Acc.
BERT _{BASE}	5:38	0.549
BERT _{LARGE}	12:38	0.562
ALBERT _{BASE}	3:16	0.490
DistilBERT _{BASE}	2:54	0.532
RoBERTa _{LARGE}	N/A	0.602

Table 3: Experiment results for classification task on SST-5 root nodes

In []: